| Character | Effect |
|---|---|
| c | Used to output the category of the logging event. The category conversion specifier can be optionally followed by *precision specifier*, that is a decimal |
| | If a precision specifier is given, then only the corresponding number of right most components of the category name will be printed. By default the category name is |
| | For example, for the category name "a.b.c" the pattern %c{2} will output "b.c". |
| C | Used to output the fully qualified class name of the caller issuing the logging request. This conversion specifier can be optionally followed by *precision* |
| | If a precision specifier is given, then only the corresponding number of right most components of the class name will be printed. By default the class name is output |
| | For example, for the class name "org.apache.xyz.SomeClass", the |
| | WARNING Generating the caller class information is slow. Thus, use should be |
| d | Used to output the date of the logging event. The date conversion specifier may be followed by a *date format specifier* enclosed between braces. For example, %d{HH:mm:ss,SSS} or %d{dd MMM yyyy HH:mm:ss,SSS}. If no date |
| | The date format specifier admits the same syntax as the time pattern string of the SimpleDateFormat. Although part of the standard JDK, the performance |
| | For better results it is recommended to use the log4j date formatters. These can be specified using one of the strings "ABSOLUTE", "DATE" and "ISO8601" for specifying AbsoluteTimeDateFormat, DateTimeDateFormat and |
| | These dedicated date formatters perform significantly better |
| F | Used to output the file name where the logging request was issued. |
| | WARNING Generating caller location information is extremely slow and should |
| l | Used to output location information of the caller which generated the logging |
| | The location information depends on the JVM implementation but usually consists of the fully qualified name of the calling method followed by the callers |
| | The location information can be very useful. However, its generation is *extremely* slow and should be avoided unless execution speed is not an issue. |
| L | Used to output the line number from where the logging request was issued. |
| | WARNING Generating caller location information is extremely slow and should |
| m | Used to output the application supplied message associated with the logging |
| M | Used to output the method name where the logging request was issued. |
| | WARNING Generating caller location information is extremely slow and should |
| n | Outputs the platform dependent line separator character or characters. |
| | This conversion character offers practically the same performance as using non-portable line separator strings such as "\n", or "\r\n". Thus, it is the preferred way |
| p | Used to output the priority of the logging event. |
| r | Used to output the number of milliseconds elapsed from the construction of the |
| t | Used to output the name of the thread that generated the logging event. |
| x | Used to output the NDC (nested diagnostic context) associated with the thread |
| X | Used to output the MDC (mapped diagnostic context) associated with the thread that generated the logging event. The X conversion character *must* be followed by the key for the map placed between braces, as |
| | See MDC class for more details. |
| % | The sequence %% outputs a single percent sign. |